

Monte-Carlo Planning: Introduction and Bandit Basics

Alan Fern

Large Worlds

- We have considered basic model-based planning algorithms
- **Model-based planning**: assumes MDP model is available
 - ▶ Methods we learned so far are at least poly-time in the number of states and actions
 - ▶ Difficult to apply to large state and action spaces (though this is a rich research area)
- We will consider various methods for overcoming this issue

Approaches for Large Worlds

- **Planning with compact MDP representations**
 1. Define a language for **compactly** describing an MDP
 - MDP is exponentially larger than description
 - E.g. via Dynamic Bayesian Networks
 2. Design a planning algorithm that directly works with that language
- Scalability is still an issue
- Can be difficult to encode the problem you care about in a given language
- Study in last part of course


Approaches for Large Worlds

- **Reinforcement learning w/ function approx.**
 1. Have a learning agent directly interact with environment
 2. Learn a compact description of policy or value function
- Often works quite well for large problems
- Doesn't fully exploit a simulator of the environment when available
- We will study reinforcement learning later in the course

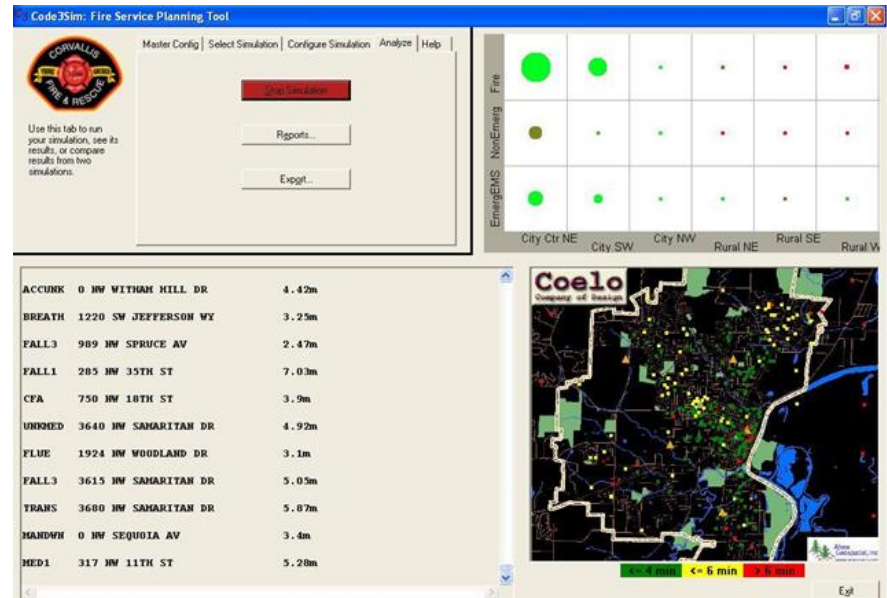
Approaches for Large Worlds: Monte-Carlo Planning

- Often a **simulator** of a planning domain is available or can be learned/estimated from data

Klondike Solitaire




Fire & Emergency Response



Large Worlds: Monte-Carlo Approach

- Often a **simulator** of a planning domain is available or can be learned from data
- **Monte-Carlo Planning**: compute a good policy for an MDP by interacting with an MDP simulator



Example Domains with Simulators

- Traffic simulators
- Robotics simulators
- Military campaign simulators
- Computer network simulators
- Emergency planning simulators
 - ▲ large-scale disaster and municipal
- Forest Fire Simulator
- Board games / Video games
 - ▲ Go / RTS

In many cases Monte-Carlo techniques yield state-of-the-art performance. Even in domains where exact MDP models are available.

MDP: Simulation-Based Representation

- A simulation-based representation gives: S , A , R , T , I :
 - ▶ finite state set S ($|S|=n$ and is generally very large)
 - ▶ finite action set A ($|A|=m$ and will assume is of reasonable size)
- $|S|$ is too large to provide a matrix representation of R , T , and I
(see next slide for I)
- A simulation based representation provides us with callable functions for R , T , and I .
 - ▶ Think of these as any other library function that you might call
- Our planning algorithms will operate by repeatedly calling those functions in an intelligent way

MDP: Simulation-Based Representation

- A simulation-based representation gives: S , A , R , T , I :
 - ▲ finite state set S ($|S|=n$ and is generally very large)
 - ▲ finite action set A ($|A|=m$ and will assume is of reasonable size)
 - ▲ Stochastic, real-valued, bounded reward function $R(s,a) = r$
 - Stochastically returns a reward r given input s and a
(note: here rewards can depend on actions and can be stochastic)
 - ▲ Stochastic transition function $T(s,a) = s'$ (i.e. a simulator)
 - Stochastically returns a state s' given input s and a
 - Probability of returning s' is dictated by $\Pr(s' \mid s,a)$ of MDP
 - ▲ Stochastic initial state function I .
 - Stochastically returns a state according to an initial state distribution


These stochastic functions can be implemented in any language!

Monte-Carlo Planning Outline

- Single State Case (multi-armed bandits)
 - ▲ A basic tool for other algorithms
- Monte-Carlo Policy Improvement
 - ▲ Policy rollout
 - ▲ Policy Switching
 - ▲ Approximate Policy Iteration
- Monte-Carlo Tree Search
 - ▲ Sparse Sampling
 - ▲ UCT and variants

Single State Monte-Carlo Planning


- Suppose MDP has a single state and k actions
 - ▲ Can sample rewards of actions using calls to simulator
 - ▲ Sampling action a is like pulling slot machine arm with random payoff function $R(s,a)$



Multi-Armed Bandit Problem

Single State Monte-Carlo Planning


- Bandit problems arise in many situations
 - ▶ Clinical trials (arms correspond to treatments)
 - ▶ Ad placement (arms correspond to ad selections)



Multi-Armed Bandit Problem

Single State Monte-Carlo Planning

- We will consider three possible bandit objectives
 - ▶ **PAC Objective:** find a near optimal arm w/ high probability
 - ▶ **Cumulative Regret:** achieve near optimal cumulative reward over lifetime of pulling (in expectation)
 - ▶ **Simple Regret:** quickly identify arm with high reward (in expectation)



Multi-Armed Bandit Problem

Multi-Armed Bandits

- Bandit algorithms are not just useful as components for multi-state Monte-Carlo planning
- Pure bandit problems arise in many applications
- Applicable whenever:
 - ▶ We have a set of independent options with unknown utilities
 - ▶ There is a cost for sampling options or a limit on total samples
 - ▶ Want to find the best option or maximize utility of our samples

Multi-Armed Bandits: Examples

- **Clinical Trials**


- ▶ Arms = possible treatments
- ▶ Arm Pulls = application of treatment to individual
- ▶ Rewards = outcome of treatment
- ▶ Objective = maximize cumulative reward = maximize benefit to trial population (or find best treatment quickly)

- **Online Advertising**

- ▶ Arms = different ads/ad-types for a web page
- ▶ Arm Pulls = displaying an ad upon a page access
- ▶ Rewards = click through
- ▶ Objective = maximize cumulative reward = maximum clicks (or find best add quickly)

PAC Bandit Objective: Informal

- **Probably Approximately Correct (PAC)**
 - ▶ Select an arm that **probably** (w/ high probability) has **approximately** the best expected reward
 - ▶ Design an algorithm that uses as few simulator calls (or pulls) as possible to guarantee this



Multi-Armed Bandit Problem

PAC Bandit Algorithms

- Let k be the number of arms, R_{max} be an upper bound on reward, and $R^* = \max_i E[R(s, a_i)]$ (i.e. R^* is the best arm reward in expectation)

Definition (Efficient PAC Bandit Algorithm): An algorithm ALG is an efficient PAC bandit algorithm iff for any multi-armed bandit problem, for any $0 < \delta < 1$ and any $0 < \varepsilon < 1$ (these are inputs to ALG), ALG pulls a number of arms that is **polynomial in $1/\varepsilon$, $1/\delta$, k , and R_{max}** and returns an arm index j such that with probability at least $1 - \delta$


$$R^* - E[R(s, a_j)] \leq \varepsilon$$

- Such an algorithm is efficient in terms of # of arm pulls, and is probably (with probability $1 - \delta$) approximately correct (picks an arm with expected reward within ε of optimal).

UniformBandit Algorithm

Even-Dar, E., Mannor, S., & Mansour, Y. (2002). PAC bounds for multi-armed bandit and Markov decision processes. In *Computational Learning Theory*

1. Pull each arm w times (uniform pulling).
2. Return arm with best average reward.



Can we make this an efficient PAC bandit algorithm?

Aside: Additive Chernoff Bound

- Let R be a random variable with maximum absolute value Z .
An let r_i $i=1, \dots, w$ be i.i.d. samples of R
- The **Chernoff bound** gives a bound on the probability that the average of the r_i are far from $E[R]$

Chernoff
Bound

$$\Pr\left(\left|E[R] - \frac{1}{w} \sum_{i=1}^w r_i\right| \geq \varepsilon\right) \leq \exp\left(-\left(\frac{\varepsilon}{Z}\right)^2 w\right)$$

Equivalent Statement:

With probability at least $1 - \delta$ we have that,

$$\left|E[R] - \frac{1}{w} \sum_{i=1}^w r_i\right| \leq Z \sqrt{\frac{1}{w} \ln \frac{1}{\delta}}$$

Aside: Coin Flip Example

- Suppose we have a coin with probability of heads equal to p .
- Let X be a random variable where $X=1$ if the coin flip gives heads and zero otherwise. (so Z from bound is 1)

$$E[X] = 1 \cdot p + 0 \cdot (1-p) = p$$


- After flipping a coin w times we can estimate the heads prob. by average of x_i .
- The Chernoff bound tells us that this estimate converges exponentially fast to the true mean (coin bias) p .

$$\Pr\left(\left|p - \frac{1}{w} \sum_{i=1}^w x_i\right| \geq \varepsilon\right) \leq \exp(-\varepsilon^2 w)$$

UniformBandit Algorithm

Even-Dar, E., Mannor, S., & Mansour, Y. (2002). PAC bounds for multi-armed bandit and Markov decision processes. In *Computational Learning Theory*

1. Pull each arm w times (uniform pulling).
2. Return arm with best average reward.



Can we make this an efficient PAC bandit algorithm?

UniformBandit PAC Bound

- For a single bandit arm the Chernoff bound says:

With probability at least $1 - \delta'$ we have that,

$$\left| E[R(s, a_i)] - \frac{1}{w} \sum_{j=1}^w r_{ij} \right| \leq R_{\max} \sqrt{\frac{1}{w} \ln \frac{1}{\delta'}}$$

- Bounding the error by ϵ gives:

$$R_{\max} \sqrt{\frac{1}{w} \ln \frac{1}{\delta'}} \leq \epsilon \quad \text{or equivalently} \quad w \geq \left(\frac{R_{\max}}{\epsilon} \right)^2 \ln \frac{1}{\delta'}$$

- Thus, using this many samples for a single arm will guarantee an ϵ -accurate estimate with probability at least $1 - \delta'$

UniformBandit PAC Bound

- So we see that with $w \geq \left(\frac{R_{\max}}{\epsilon}\right)^2 \ln \frac{1}{\delta'}$ samples per arm, there is no more than a δ' probability that an individual arm's estimate will **not** be ϵ -accurate
 - ▲ But we want to bound the probability of any arm being inaccurate

The **union bound** says that for k events, the probability that at least one event occurs is bounded by the sum of individual probabilities

$$\Pr(A_1 \text{ or } A_2 \text{ or } \cdots \text{ or } A_k) \leq \sum_{i=1}^k \Pr(A_i)$$

- Using the above # samples per arm and the union bound (with events being “arm i is not ϵ -accurate”) there is no more than $k\delta'$ probability of any arm not being ϵ -accurate
- Setting $\delta' = \frac{\delta}{k}$ all arms are ϵ -accurate with prob. at least $1 - \delta$

UniformBandit PAC Bound


Putting everything together we get:

$$\text{If } w \geq \left(\frac{R_{\max}}{\varepsilon} \right)^2 \ln \frac{k}{\delta} \text{ then for all arms simultaneously}$$
$$\left| E[R(s, a_i)] - \frac{1}{w} \sum_{j=1}^w r_{ij} \right| \leq \varepsilon$$

with probability at least $1 - \delta$


- That is, estimates of all actions are ϵ -accurate with probability at least $1 - \delta$
- Thus selecting estimate with highest value is approximately optimal with high probability, or PAC

Simulator Calls for UniformBandit



- Total simulator calls for PAC: $k \cdot w = \left(\frac{R_{\max}}{\varepsilon} \right)^2 k \ln \frac{k}{\delta}$
 - So we have an **efficient** PAC algorithm
 - Can we do better than this?

Non-Uniform Sampling



- If an arm is really bad, we should be able to eliminate it from consideration early on
- **Idea:** try to allocate more pulls to arms that appear more promising

Median Elimination Algorithm

Even-Dar, E., Mannor, S., & Mansour, Y. (2002). PAC bounds for multi-armed bandit and Markov decision processes. In *Computational Learning Theory*

Median Elimination

A = set of all arms

For $i = 1$ to

 Pull each arm in A w_i times

m = median of the average rewards of the arms in A

$A = A - \{\text{arms with average reward less than } m\}$

 If $|A| = 1$ then return the arm in A

Eliminates half of the arms each round.

How to set the w_i to get PAC guarantee?

Median Elimination (proof not covered)

- Theoretical values used by Median Elimination:

$$w_i = \frac{4}{\epsilon_i^2} \ln \frac{3}{\delta_i} \quad \epsilon_i = \left(\frac{3}{4}\right)^{i-1} \cdot \frac{\epsilon}{4} \quad \delta_i = \frac{\delta}{2^i}$$

Theorem: Median Elimination is a PAC algorithm and uses a number of pulls that is at most $O\left(\frac{k}{\epsilon^2} \ln \frac{1}{\delta}\right)$


Compare to $O\left(\frac{k}{\epsilon^2} \ln \frac{k}{\delta}\right)$ for UniformBandit

PAC Summary

- Median Elimination uses $O(\log(k))$ fewer pulls than Uniform
 - Known to be asymptotically optimal (no PAC algorithm can use fewer pulls in worst case)
- PAC objective is sometimes awkward in practice
 - Sometimes we don't know how many pulls we will have
 - Sometimes we can't control how many pulls we get
 - Selecting ϵ and δ can be quite arbitrary
- Cumulative & simple regret partly address this

Cumulative Regret Objective

- **Problem:** find arm-pulling strategy such that the expected total reward at time n is close to the best possible (one pull per time step)
 - ▶ Optimal (in expectation) is to pull optimal arm n times
 - ▶ UniformBandit is poor choice --- waste time on bad arms
 - ▶ Must balance **exploring** machines to find good payoffs and **exploiting** current knowledge



Cumulative Regret Objective

- Theoretical results are often about “expected cumulative regret” of an arm pulling strategy.
- **Protocol:** At time step n the algorithm picks an arm a_n based on what it has seen so far and receives reward r_n (a_n and r_n are random variables).
- **Expected Cumulative Regret ($E[Reg_n]$):** difference between optimal expected cumulative reward and expected cumulative reward of our strategy at time n

$$E[Reg_n] = n \cdot R^* - \sum_{i=1}^n E[r_i]$$

UCB Algorithm for Minimizing Cumulative Regret

Auer, P., Cesa-Bianchi, N., & Fischer, P. (2002). Finite-time analysis of the multiarmed bandit problem. *Machine learning*, 47(2), 235-256.

- $Q(a)$: average reward for trying action a (in our single state s) so far
- $n(a)$: number of pulls of arm a so far
- Action choice by UCB after n pulls:

$$a_n = \arg \max_a Q(a) + \sqrt{\frac{2 \ln n}{n(a)}}$$

- Assumes rewards in $[0,1]$. We can always normalize if we know max value.

UCB: Bounded Sub-Optimality

$$a_n = \arg \max_a Q(a) + \sqrt{\frac{2 \ln n}{n(a)}}$$

Value Term:

favors actions that looked
good historically

Exploration Term:

actions get an exploration
bonus that grows with $\ln(n)$

Expected number of pulls of sub-optimal arm **a** is bounded by:

$$\frac{8}{\Delta_a^2} \ln n$$

where Δ_a is the sub-optimality of arm **a**

Doesn't waste much time on sub-optimal arms, unlike uniform!

UCB Performance Guarantee

[Auer, Cesa-Bianchi, & Fischer, 2002]

Theorem: The expected cumulative regret of UCB $E[\text{Reg}_n]$ after n arm pulls is bounded by $O(\log n)$

- Is this good?

Yes. The average per-step regret is $O\left(\frac{\log(n)}{n}\right)$

Theorem: No algorithm can achieve a better expected regret (up to constant factors)

What Else

- UCB is great when we care about cumulative regret
- But, sometimes all we care about is finding a good arm quickly
- This is similar to the PAC objective, but:
 - ▲ The PAC algorithms required precise knowledge of or control of # pulls
 - ▲ We would like to be able to stop at any time and get a good result with some guarantees on expected performance
- “Simple regret” is an appropriate objective in these cases

Simple Regret Objective

- **Protocol:** At time step n the algorithm picks an “exploration” arm a_n to pull and observes reward r_n and also picks an arm index it thinks is best j_n (a_n , j_n and r_n are random variables).
 - ▲ If interrupted at time n the algorithm returns j_n .
- **Expected Simple Regret ($E[SReg_n]$):** difference between R^* and expected reward of arm j_n selected by our strategy at time n

$$E[SReg_n] = R^* - E[R(a_{j_n})]$$

Simple Regret Objective

- What about UCB for simple regret?
 - Intuitively we might think UCB puts too much emphasis on pulling the best arm
 - After an arm starts looking good, we might be better off trying figure out if there is indeed a better arm

Theorem: The expected simple regret of UCB after n arm pulls is upper bounded by $O(n^{-c})$ for a constant c .

Seems good, but we can do much better in theory.

Incremental Uniform (or Round Robin)

Bubeck, S., Munos, R., & Stoltz, G. (2011). Pure exploration in finitely-armed and continuous-armed bandits. Theoretical Computer Science, 412(19), 1832-1852

Algorithm:

- At round n pull arm with index $(k \bmod n) + 1$
- At round n return arm (if asked) with largest average reward

Theorem: The expected simple regret of Uniform after n arm pulls is upper bounded by $O(e^{-cn})$ for a constant c .

- This bound is exponentially decreasing in n !

Compared to polynomially for UCB $O(n^{-c})$.

Can we do better?

Tolpin, D. & Shimony, S, E. (2012). MCTS Based on Simple Regret. *AAAI Conference on Artificial Intelligence*.

Algorithm ϵ -Greedy : (parameter $0 < \epsilon < 1$)

- At round n , with probability ϵ pull arm with best average reward so far, otherwise pull one of the other arms at random.
- At round n return arm (if asked) with largest average reward

Theorem: The expected simple regret of ϵ -Greedy for $\epsilon = 0.5$ after n arm pulls is upper bounded by $O(e^{-cn})$ for a constant c that is larger than the constant for Uniform (this holds for “large enough” n).

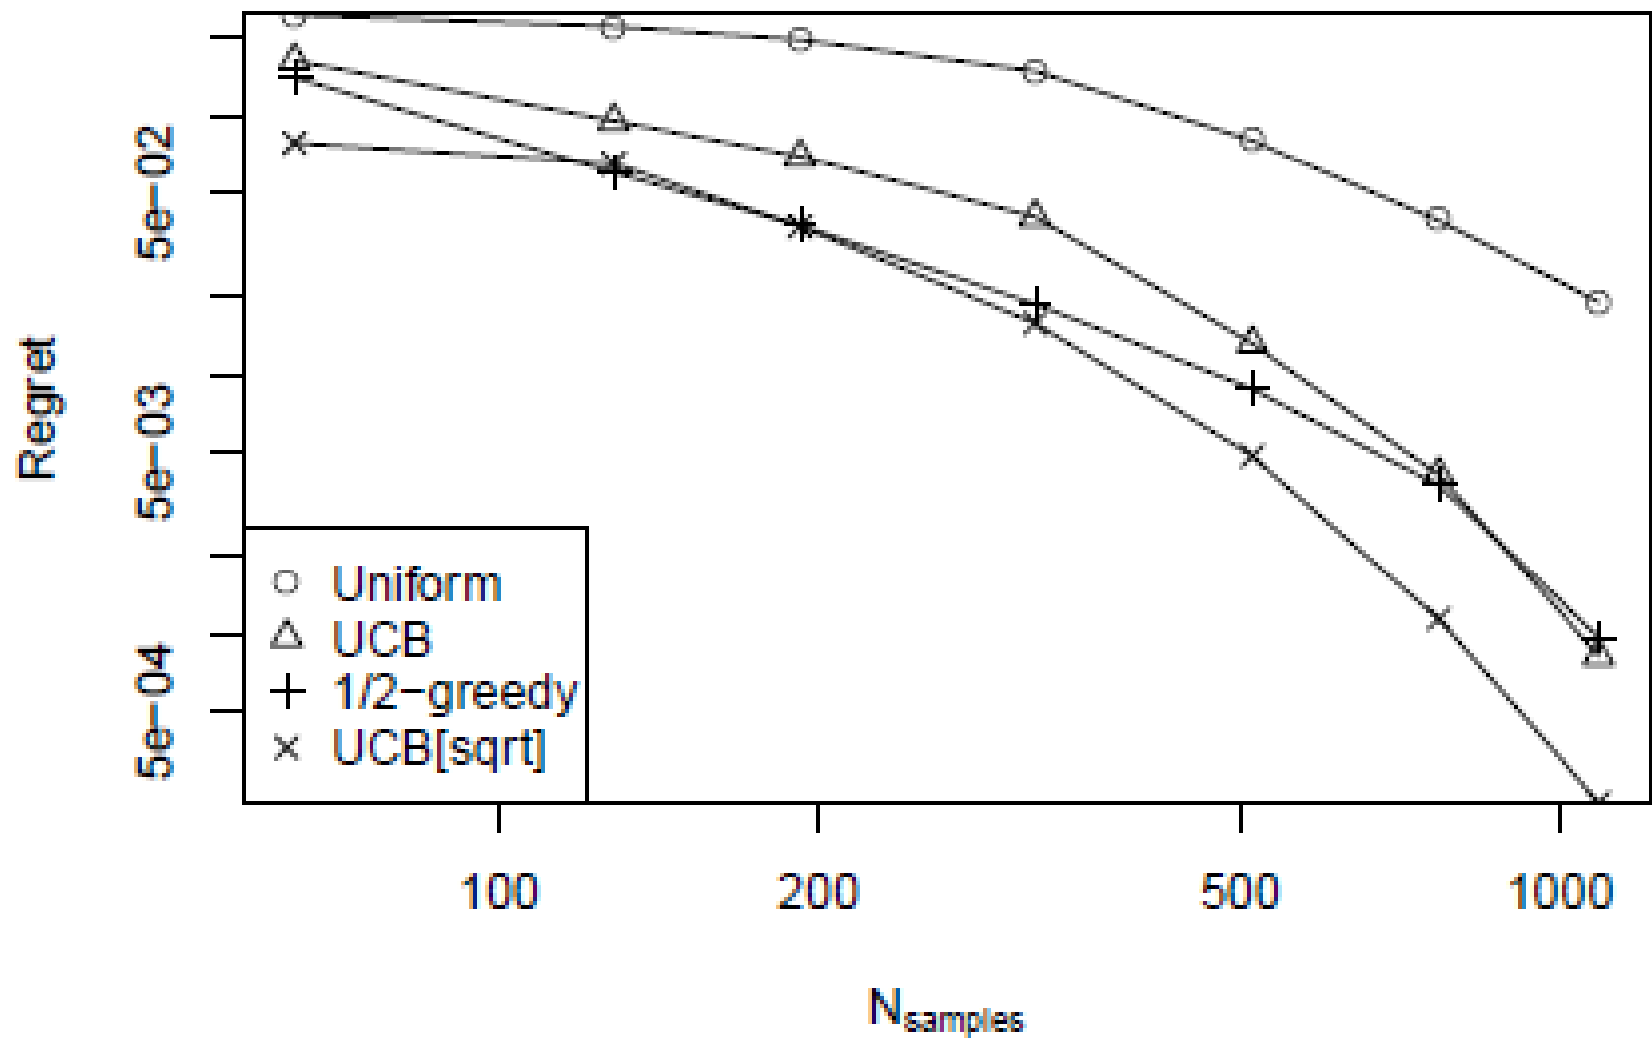
Summary of Bandits in Theory

- PAC Objective:
 - **UniformBandit** is a simple PAC algorithm
 - **MedianElimination** improves by a factor of $\log(k)$ and is optimal up to constant factors
- Cumulative Regret:
 - **Uniform** is very bad!
 - **UCB** is optimal (up to constant factors)
- Simple Regret:
 - **UCB** shown to reduce regret at polynomial rate
 - **Uniform** reduces at an exponential rate
 - **0.5-Greedy** may have even better exponential rate

Theory vs. Practice

- The established theoretical relationships among bandit algorithms have often been useful in predicting empirical relationships.
- But not always

Theory vs. Practice



b. regret vs. number of samples