

Monte-Carlo Planning: Policy Improvement

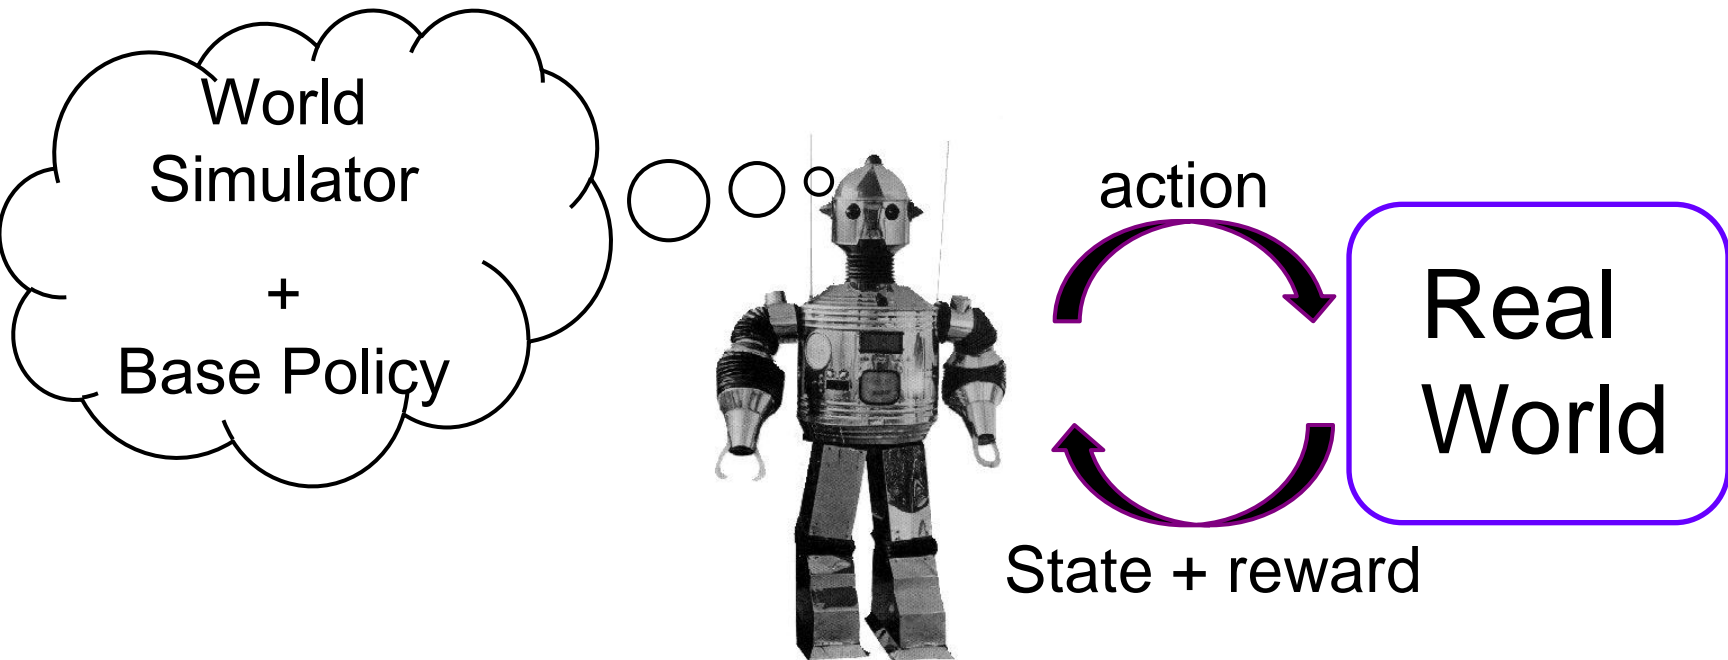
Alan Fern

Monte-Carlo Planning Outline

- Single State Case (multi-armed bandits)
 - ▲ A basic tool for other algorithms
- Monte-Carlo Policy Improvement
 - ▲ Policy rollout
 - ▲ Policy Switching
- Monte-Carlo Tree Search
 - ▲ Sparse Sampling
 - ▲ UCT and variants

Policy Improvement via Monte-Carlo

- Now consider a very large multi-state MDP.
- Suppose we have a simulator and a non-optimal policy
 - ▲ E.g. policy could be a standard heuristic or based on intuition
- Can we somehow compute an improved policy?




Recall: Policy Improvement Theorem

$$Q_{\pi}(s, a) = R(s, a) + \beta \sum_{s'} T(s, a, s') \cdot V_{\pi}(s')$$

- The Q-value function of a policy gives **expected discounted future reward of starting in state s , taking action a , and then following policy π thereafter**
- **Define:** $\pi'(s) = \arg \max_a Q_{\pi}(s, a)$
- **Theorem [Howard, 1960]:** For any non-optimal policy π the policy π' a strict improvement over π .
- Computing π' amounts to finding the action that maximizes the Q-function of π
 - ▶ Can we use the bandit idea to solve this?

Policy Improvement via Bandits



- **Idea:** define a stochastic function **SimQ(s,a,π)** that we can implement and whose expected value is $Q_\pi(s,a)$
- Then use Bandit algorithm to select (approx) best action


How to implement SimQ?

Q-value Estimation

- SimQ might be implemented by simulating the execution of action a in state s and then following π thereafter
 - ▶ But for infinite horizon problems this would never finish
 - ▶ So we will approximate via finite horizon
- The h -horizon Q-function $Q_\pi(s, a, h)$ is defined as:
expected total discounted reward of starting in state s , taking action a , and then following policy π for $h-1$ steps
- The approximation error decreases exponentially fast in h

$$\left| Q_\pi(s, a) - Q_\pi(s, a, h) \right| \leq \beta^h V_{\max} \quad V_{\max} = \frac{R_{\max}}{1 - \beta}$$

Policy Improvement via Bandits



- **Refined Idea:** define a stochastic function **$\text{SimQ}(s, a, \pi, h)$** that we can implement, whose expected value is $Q_\pi(s, a, h)$
- Use Bandit algorithm to select (approx) best action

How to implement SimQ?

Policy Improvement via Bandits

SimQ(s,a, π ,h)

$r = R(s,a)$

$s = T(s,a)$

for $i = 1$ to $h-1$

$r = r + \beta^i R(s, \pi(s))$

$s = T(s, \pi(s))$

Return r

} simulate a in s

} simulate $h-1$ steps
of policy

- Simply simulate taking a in s and following policy for $h-1$ steps, returning discounted sum of rewards
- Expected value of SimQ(s,a, π ,h) is $Q_{\pi}(s,a,h)$ which can be made arbitrarily close to $Q_{\pi}(s,a)$ by increasing h

Policy Improvement via Bandits

$\text{SimQ}(s, a, \pi, h)$

$r = R(s, a)$

$s = T(s, a)$

for $i = 1$ to $h-1$

$r = r + \beta^i R(s, \pi(s))$


$s = T(s, \pi(s))$

Return r


} simulate a in s

} simulate $h-1$ steps
of policy

Trajectory under π



Policy Improvement via Bandits



- **Refined Idea:** define a stochastic function **$\text{SimQ}(s, a, \pi, h)$** that we can implement, whose expected value is $Q_\pi(s, a, h)$
- Use Bandit algorithm to select (approx) best action

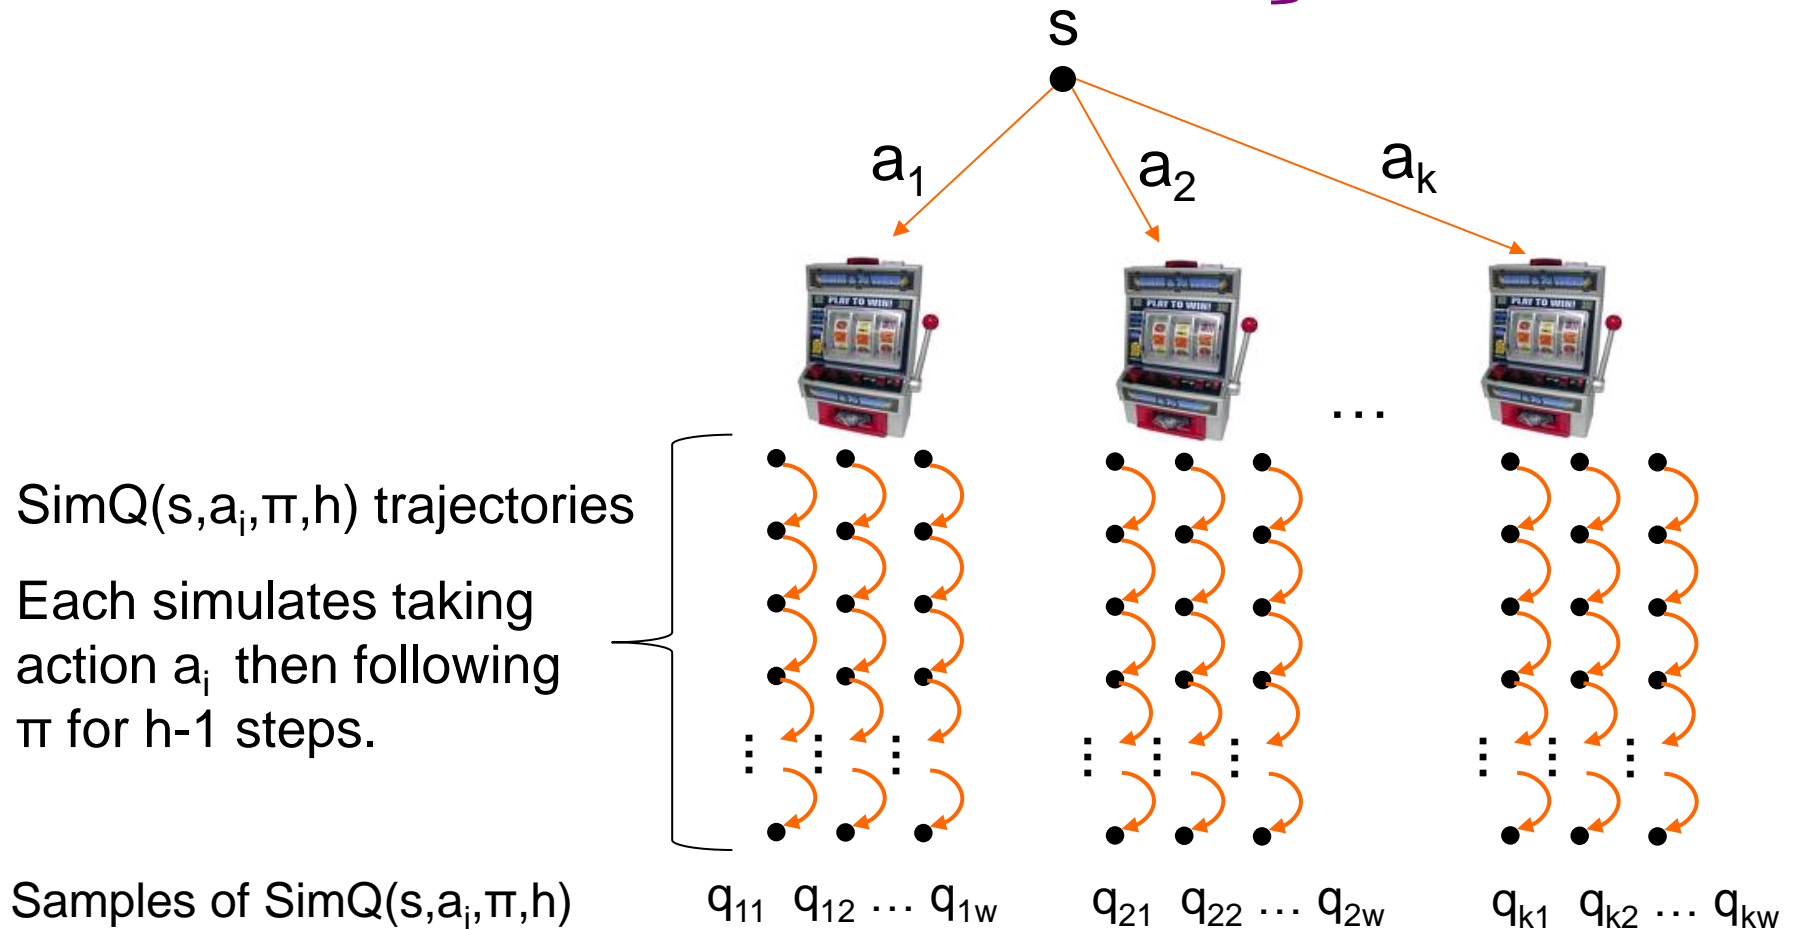
Which bandit objective/algorithm to use?

Traditional Approach: Policy Rollout


UniformRollout[π, h, w](s)

1. For each a_i run $\text{SimQ}(s, a_i, \pi, h)$ w times
2. Return action with best average of SimQ results

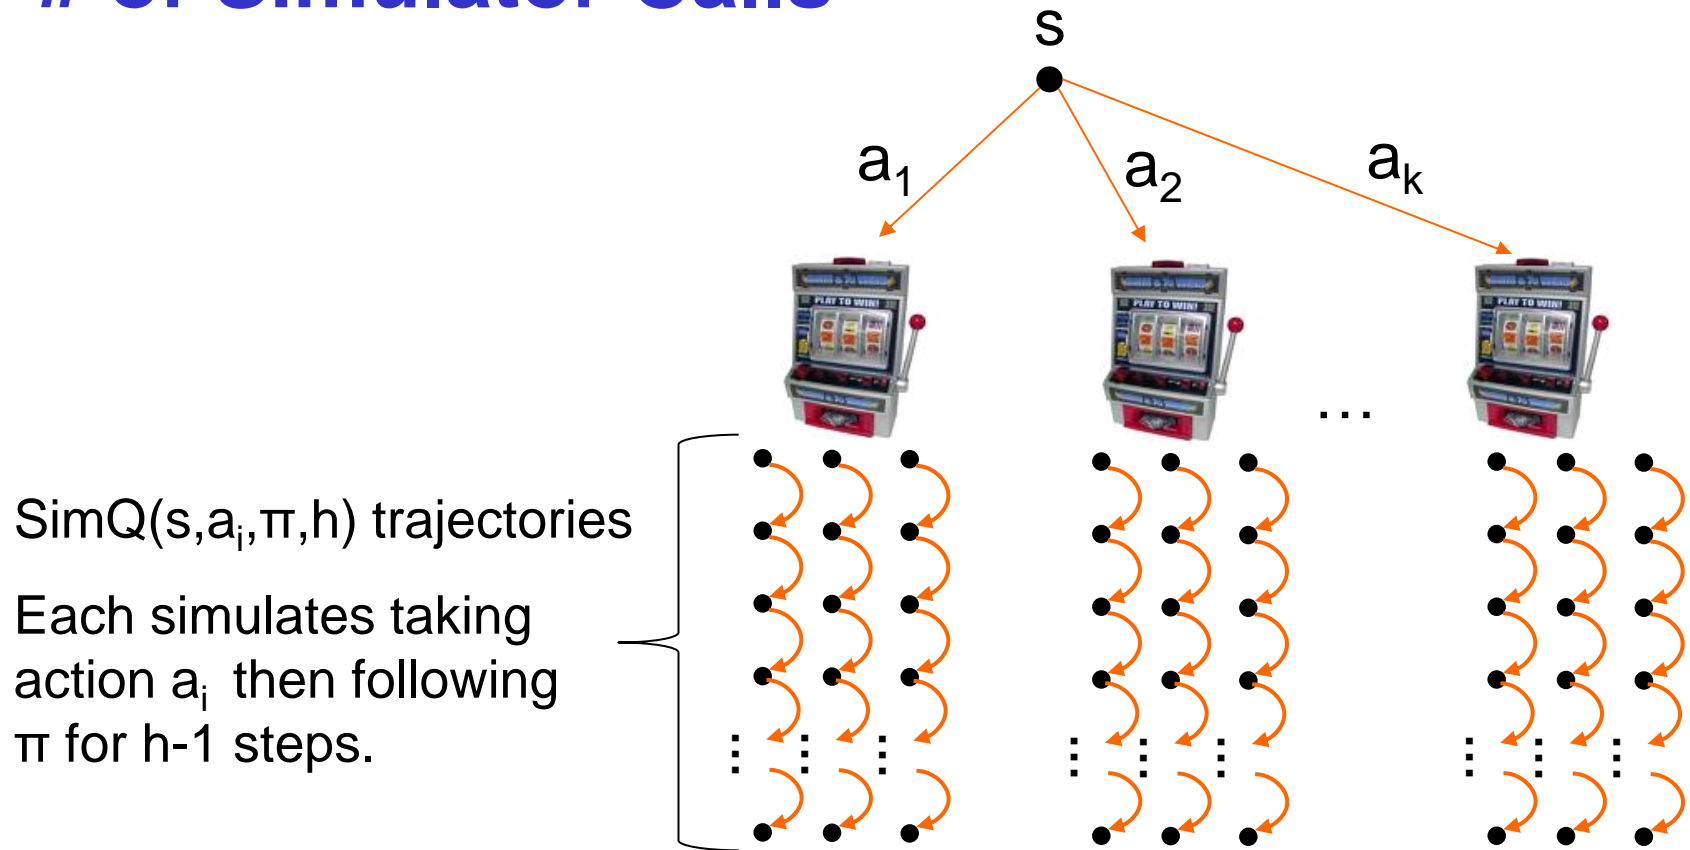
UniformBandit
for PAC objective



Executing Rollout in Real World



Uniform Policy Rollout: # of Simulator Calls



- For each action w calls to SimQ, each using h sim calls
- Total of khw calls to the simulator

Uniform Policy Rollout: PAC Guarantee

- Let a^* be the action that maximizes the true Q-function $Q_\pi(s, a)$.
- Let a' be the action returned by $\text{UniformRollout}[\pi, h, w](s)$.
- Putting the PAC bandit result together with the finite horizon approximation we can derive the following:

$$\text{If } w \geq \left(\frac{R_{\max}}{\varepsilon} \right)^2 \ln \frac{k}{\delta} \text{ then with probability at least } 1 - \delta$$
$$|Q_\pi(s, a^*) - Q_\pi(s, a')| \leq \varepsilon + \beta^h V_{\max}$$

But does this guarantee that the value of $\text{UniformRollout}[\pi, h, w](s)$ will be close to the value of π' ?

Policy Rollout: Quality

- How good is $\text{UniformRollout}[\pi, h, w]$ compared to π' ?
- **Bad News.** In general for a fixed h and w there is always an MDP such that the quality of the rollout policy is arbitrarily worse than π' .
- The example MDP is somewhat involved, but shows that even small error in Q-value estimates can lead to large performance gaps compared to π'
 - ▲ But this result is quite pathological


Policy Rollout: Quality

- How good is $\text{UniformRollout}[\pi, h, w]$ compared to π' ?
- **Good News.** If we make an assumption about the MDP, then it is possible to select h and w so that the rollout quality is close to π' .
 - ▲ This is a bit involved.
 - ▲ Assume a lower bound on the difference between the best Q-value and the second best Q-value
- **More Good News.** It is possible to select h and w so that $\text{Rollout}[\pi, h, w]$ is (approximately) no worse than π for any MDP
 - ▲ So at least rollout won't hurt compared to the base policy
 - ▲ At the same time it has the potential to significantly help

Non-Uniform Policy Rollout

- Should we consider minimizing cumulative regret?

No! We really only care about finding an (approx) best arm.




Non-Uniform Policy Rollout

PAC Setting: use **MedianElimination**

(parameterized by ϵ and δ instead of w)

- Often we are given a budget on number of samples (i.e. time per decision).
- MedianElimination not applicable.




Non-Uniform Policy Rollout

Simple Regret: use ϵ -Greedy

(parameterized by budget n on # of pulls)

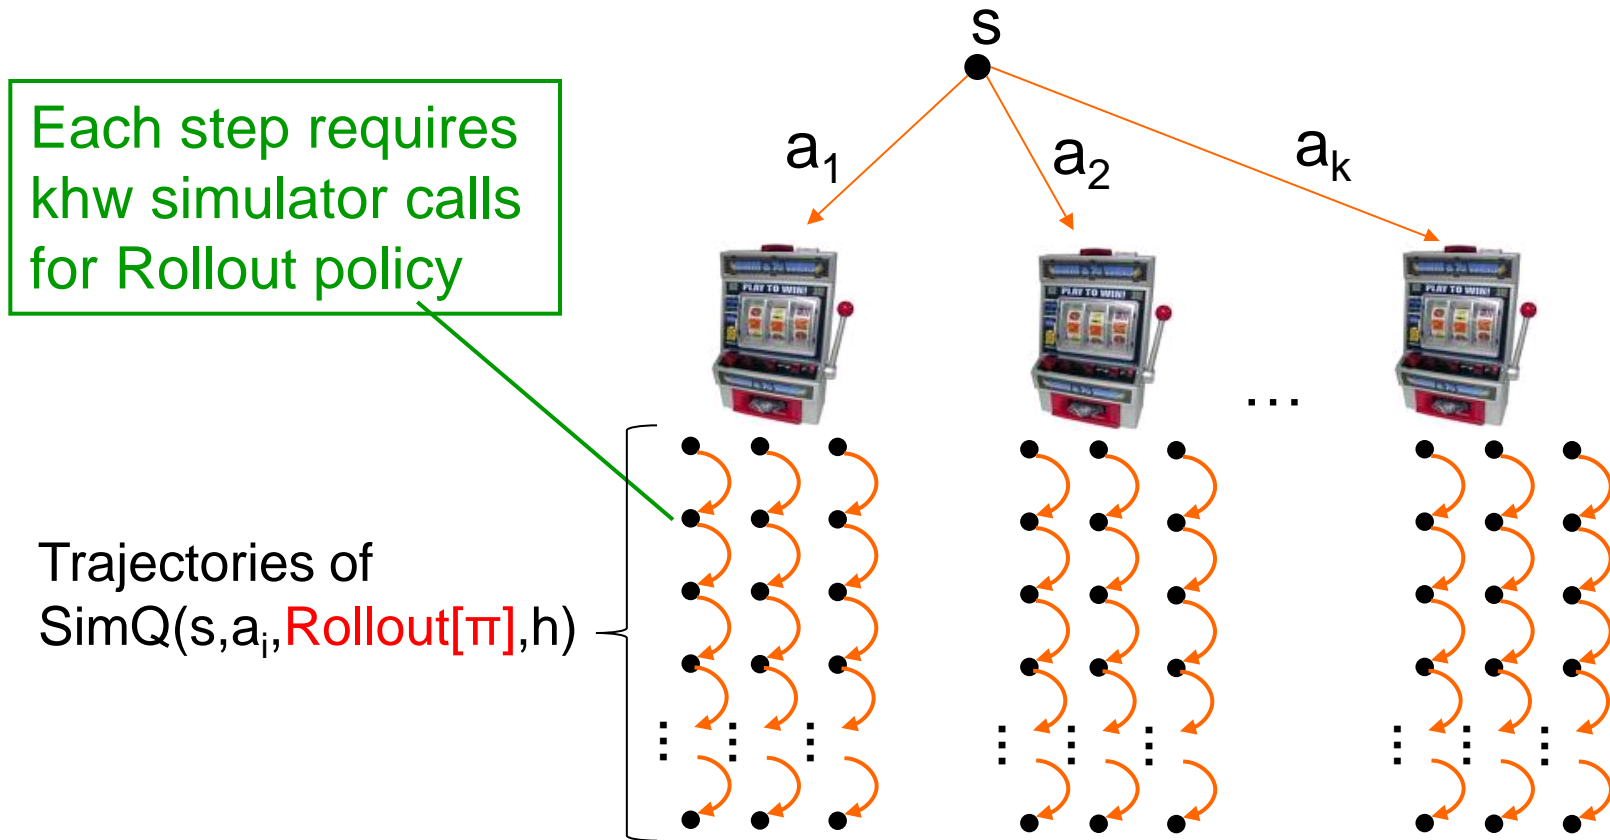
- Call this ϵ -Rollout $[\pi, h, n]$
- n is number of samples per step
- For $\epsilon = 0.5$ we might expect it to be better than UniformRollout for same # of total samples.



Multi-Stage Rollout

- In what follows we will use the notation **Rollout[π]** to refer to either UniformRollout[π, h, w] or ϵ -Rollout[π, h, n].
- A single call to Rollout[π](s) approximates one iteration of policy iteration initialized at policy π
 - ▶ But only computes the action for state s rather than all states (as done by full policy iteration)!
- We can use more computation time to approximate multiple iterations of policy iteration via nesting calls to Rollout
- Gives a way to use more time in order to improve performance

Multi-Stage Rollout



- Two stage: compute rollout policy of “rollout policy of π ”
- Requires $(khw)^2$ calls to the simulator for 2 stages
- In general exponential in the number of stages

Rollout Summary

- We often are able to write simple, mediocre policies
 - ▲ Network routing policy
 - ▲ Policy for card game of Hearts
 - ▲ Policy for game of Backgammon
 - ▲ Solitaire playing policy
- Policy rollout is a general and easy way to improve upon such policies given a simulator
- Often observe substantial improvement, e.g.
 - ▲ Compiler instruction scheduling
 - ▲ Backgammon
 - ▲ Network routing
 - ▲ Combinatorial optimization
 - ▲ Game of GO
 - ▲ Solitaire

Example: Rollout for Solitaire [Yan et al. NIPS'04]

Player	Success Rate	Time/Game
Human Expert	36.6%	20 min
(naïve) Base Policy	13.05%	0.021 sec
1 rollout	31.20%	0.67 sec
2 rollout	47.6%	7.13 sec
3 rollout	56.83%	1.5 min
4 rollout	60.51%	18 min
5 rollout	70.20%	1 hour 45 min

- Multiple levels of rollout can payoff but is expensive

Monte-Carlo Planning Outline

- Single State Case (multi-armed bandits)
 - ▲ A basic tool for other algorithms
- Monte-Carlo Policy Improvement
 - ▲ Policy rollout
 - ▲ Policy Switching
- Monte-Carlo Tree Search
 - ▲ Sparse Sampling
 - ▲ UCT and variants


Another Useful Technique: Policy Switching

- Sometimes policy rollout can be too expensive when the number of actions is large (time scales linearly with number of actions)
- Sometimes we have multiple base policies and it is hard to pick just one to use for rollout.
- Policy switching helps deal with both of these issues.

Another Useful Technique: Policy Switching

- Suppose you have a set of base policies $\{\pi_1, \pi_2, \dots, \pi_M\}$
- Also suppose that the best policy to use can depend on the specific state of the system and we don't know how to select.
- Policy switching is a simple way to select which policy to use at a given step via a simulator

Another Useful Technique: Policy Switching

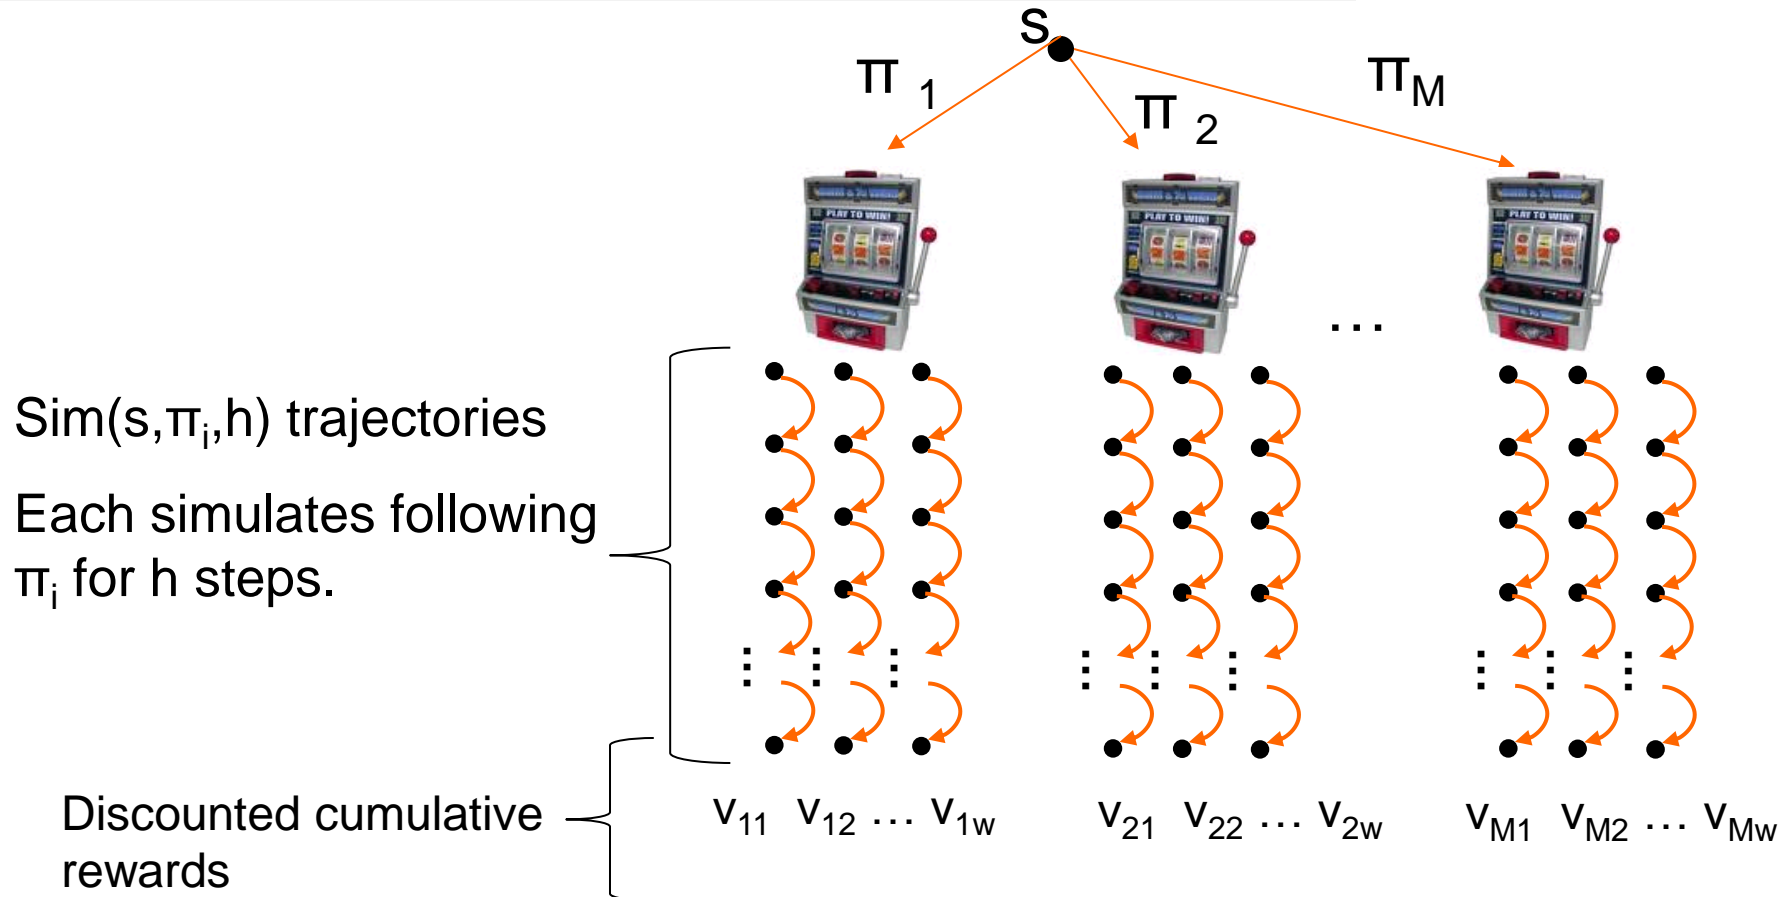


- The stochastic function **$\text{Sim}(s, \pi, h)$** simply samples the h -horizon value of π starting in state s
- Implement by simply simulating π starting in s for h steps and returning discounted total reward
- Use Bandit algorithm to select best policy and then select action chosen by that policy


Uniform Policy Switching

UniformPolicySwitch $[\{\pi_1, \pi_2, \dots, \pi_M\}, h, w](s)$

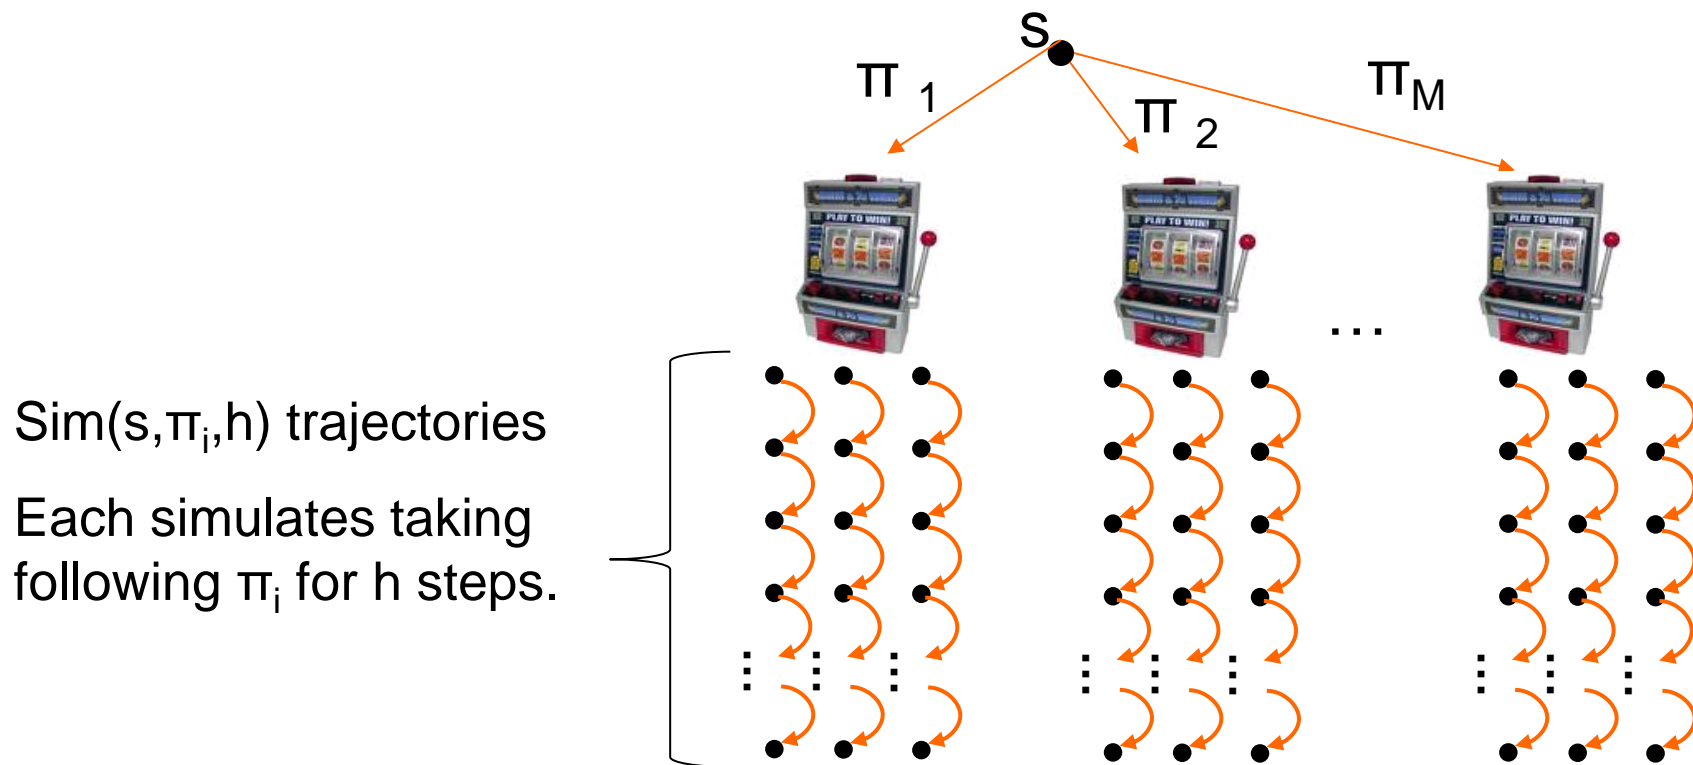
1. For each π_i run $\text{Sim}(s, \pi_i, h)$ w times
2. Let i^* be index of policy with best average result
3. Return action $\pi_{i^*}(s)$



Executing Policy Switching in Real World



Uniform Policy Switching: Simulator Calls




- For each policy use w calls to Sim , each using h simulator calls
- Total of Mhw calls to the simulator
- Does not depend on number of actions!

ϵ -Greedy Policy Switching

- Similar to rollout we can have a non-uniform version that takes a total number of trajectories n as an argument

ϵ -PolicySwitch $[\{\pi_1, \dots, \pi_M\}, h, n]$

Use ϵ -Greedy as the bandit algorithm for n pulls and return best arm/policy.



Policy Switching: Quality

- Let π_{ps} denote the ideal switching policy
 - ▲ Always pick the best policy index at any state

Theorem: For any state s , $\max_i V_{\pi_i}(s) \leq V_{\pi_{ps}}(s)$.

- The value of the switching policy is at least as good as the best single policy in the set
 - ▲ It will often perform better than any single policy in set.
 - ▲ For non-ideal case, where bandit algorithm only picks approximately the best arm we can add an error term to the bound.

Proof

Theorem: For any state s , $\max_i V_{\pi_i}(s) \leq V_{\pi_{ps}}(s)$.

We'll use the following property.

Proposition: For any policy π and value function V ,
if $V \leq B_{\pi}[V]$, then $V \leq V_{\pi}$

Recall $B_{\pi}[V](s) = R(s) + \sum_{s'} T(s, \pi(s), s') \cdot V(s')$
is the restricted Bellman backup.

So all we need to do is prove that $\max_i V_{\pi_i} \leq B_{\pi_{ps}} \left[\max_i V_{\pi_i} \right]$
since this will imply that $\max_i V_{\pi_i} \leq V_{\pi_{ps}}$ as desired.

Proof *(to simply notation and without loss of generality, assume rewards only depend on state and are deterministic)*

Prove that $\max_i V_{\pi_i} \leq B_{\pi_{ps}} \left[\max_i V_{\pi_i} \right]$

Let i^* be the index of the best policy in state s .

$$\begin{aligned} B_{\pi_{ps}} \left[\max_i V_{\pi_i} \right] (s) &= R(s) + \sum_{s'} T(s, \pi_{ps}(s), s') \cdot \max_i V_{\pi_i}(s') \\ &\geq R(s) + \max_i \sum_{s'} T(s, \pi_{i^*}(s), s') \cdot V_{\pi_i}(s') \\ &= \max_i \left[R(s) + \sum_{s'} T(s, \pi_{i^*}(s), s') \cdot V_{\pi_i}(s') \right] \\ &\geq \max_i \left[R(s) + \sum_{s'} T(s, \pi_i(s), s') \cdot V_{\pi_i}(s') \right] \\ &= \max_i V_{\pi_i}(s) \end{aligned}$$

Policy Switching Summary

- Easy way to produce an improved policy from a set of existing policies.
 - ▲ Will not do any worse than the best policy in your set.
- Complexity does not depend on number of actions.
 - ▲ So can be practical even when action space is huge, unlike policy rollout.
- Can combine with rollout for further improvement
 - ▲ Just apply rollout to the switching policy.