

Reinforcement Learning

Introduction & Passive Learning

Alan Fern

* Based in part on slides by Daniel Weld

So far

- Given an MDP model we know how to find optimal policies (for moderately-sized MDPs)
 - ▲ Value Iteration or Policy Iteration
- Given just a simulator of an MDP we know how to select actions
 - ▲ Monte-Carlo Planning
- What if we don't have a model or simulator?
 - ▲ Like when we were babies . . .
 - ▲ Like in many real-world applications
 - ▲ All we can do is wander around the world observing what happens, getting rewarded and punished
- Enters reinforcement learning

Reinforcement Learning

- No knowledge of environment
 - ▲ Can only act in the world and observe states and reward
- Many factors make RL difficult:
 - ▲ Actions have **non-deterministic effects**
 - Which are initially unknown
 - ▲ **Rewards / punishments** are infrequent
 - Often at the end of long sequences of actions
 - How do we determine what action(s) were really responsible for reward or punishment?
(credit assignment)
 - ▲ World is large and complex
- Nevertheless learner **must decide** what actions to take
 - ▲ We will assume the world behaves as an MDP

Pure Reinforcement Learning vs. Monte-Carlo Planning

- In pure reinforcement learning:
 - ▲ the agent begins with no knowledge
 - ▲ wanders around the world observing outcomes
- In Monte-Carlo planning
 - ▲ the agent begins with no declarative knowledge of the world
 - ▲ has an interface to a world simulator that allows observing the outcome of taking any action in any state
- The simulator gives the agent the ability to “teleport” to any state, at any time, and then apply any action
- A pure RL agent does not have the ability to teleport
 - ▲ Can only observe the outcomes that it happens to reach

Pure Reinforcement Learning vs. Monte-Carlo Planning

- MC planning is sometimes called RL with a “strong simulator”
 - ▲ I.e. a simulator where we can set the current state to any state at any moment
- Pure RL is sometimes called RL with a “weak simulator”
 - ▲ I.e. a simulator where we cannot set the state
- A strong simulator can emulate a weak simulator
 - ▲ So pure RL can be used in the MC planning framework
 - ▲ But not vice versa

Passive vs. Active learning

- Passive learning
 - ▲ The agent has a fixed policy and tries to learn the utilities of states by observing the world go by
 - ▲ Analogous to policy evaluation
 - ▲ Often serves as a component of active learning algorithms
 - ▲ Often inspires active learning algorithms
- Active learning
 - ▲ The agent attempts to find an optimal (or at least good) policy by acting in the world
 - ▲ Analogous to solving the underlying MDP, but without first being given the MDP model

Model-Based vs. Model-Free RL


- *Model based approach to RL:*
 - ▲ learn the MDP model, or an approximation of it
 - ▲ use it for policy evaluation or to find the optimal policy
- *Model free approach to RL:*
 - ▲ derive the optimal policy without explicitly learning the model
 - ▲ useful when model is difficult to represent and/or learn
- We will consider both types of approaches

Small vs. Huge MDPs

- We will first cover RL methods for small MDPs
 - ▲ MDPs where the number of states and actions is reasonably small
 - ▲ These algorithms will inspire more advanced methods
- Later we will cover algorithms for huge MDPs
 - ▲ Function Approximation Methods
 - ▲ Policy Gradient Methods
 - ▲ Least-Squares Policy Iteration

Example: Passive RL

- Suppose given a stationary policy (shown by arrows)
 - ▲ Actions can stochastically lead to unintended grid cell
- Want to determine how good it is




Objective: Value Function

3	0.812	0.868	0.918	<div>+ 1</div>
2	0.762		0.660	<div>- 1</div>
1	0.705	0.655	0.611	0.388
	1	2	3	4

Passive RL

- Estimate $V^\pi(s)$
- Not given
 - ▶ transition matrix, nor
 - ▶ reward function!



- Follow the policy for many epochs giving training sequences.

$(1,1) \rightarrow (1,2) \rightarrow (1,3) \rightarrow (1,2) \rightarrow (1,3) \rightarrow (2,3) \rightarrow (3,3) \rightarrow (3,4)$ +1

$(1,1) \rightarrow (1,2) \rightarrow (1,3) \rightarrow (2,3) \rightarrow (3,3) \rightarrow (3,2) \rightarrow (3,3) \rightarrow (3,4)$ +1

$(1,1) \rightarrow (2,1) \rightarrow (3,1) \rightarrow (3,2) \rightarrow (4,2)$ -1

- Assume that after entering +1 or -1 state the agent enters zero reward terminal state
 - ▶ So we don't bother showing those transitions

Approach 1: Direct Estimation

- Direct estimation (also called Monte Carlo)
 - ▲ Estimate $V^\pi(s)$ as average total reward of epochs containing s (calculating from s to end of epoch)
- ***Reward to go*** of a state s

the sum of the (discounted) rewards from that state until a terminal state is reached
- Key: use observed ***reward to go*** of the state as the direct evidence of the actual expected utility of that state
- Averaging the reward-to-go samples will converge to true value at state

Direct Estimation

- Converge very slowly to correct utilities values (requires a lot of sequences)
- Doesn't exploit Bellman constraints on policy values

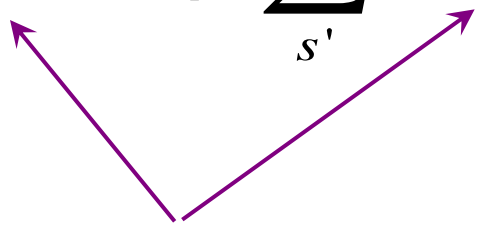
$$V^{\pi}(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^{\pi}(s')$$

- ▲ It is happy to consider value function estimates that violate this property badly.

How can we incorporate the Bellman constraints?

Approach 2: Adaptive Dynamic Programming (ADP)


- ADP is a model based approach
 - ▶ Follow the policy for awhile
 - ▶ Estimate transition model based on observations
 - ▶ Learn reward function
 - ▶ Use estimated model to compute utility of policy

$$V^{\pi}(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^{\pi}(s')$$


learned

- How can we estimate transition model $T(s, a, s')$?
 - ▶ Simply the fraction of times we see s' after taking a in state s .
 - ▶ NOTE: Can bound error with Chernoff bounds if we want

ADP learning curves



Approach 3: Temporal Difference Learning (TD)

- Can we avoid the computational expense of full DP policy evaluation?
- Can we avoid the $O(n^2)$ space requirements for storing the transition model estimate?
- Temporal Difference Learning (model free)
 - ▲ Doesn't store an estimate of entire transition function
 - ▲ Instead stores estimate of V^π , which requires only $O(n)$ space.
 - ▲ Does local, cheap updates of utility/value function on a per-action basis

Approach 3: Temporal Difference Learning (TD)

For each transition of π from s to s' , update $V^\pi(s)$ as follows:

$$V^\pi(s) \leftarrow V^\pi(s) + \alpha(R(s) + \beta V^\pi(s') - V^\pi(s))$$

updated estimate learning rate discount factor current estimates at s' and s

- Intuitively moves us closer to satisfying Bellman constraint

$$V^\pi(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^\pi(s')$$

Why?

Aside: Online Mean Estimation

- Suppose that we want to incrementally compute the mean of a sequence of numbers (x_1, x_2, x_3, \dots)
 - ▲ E.g. to estimate the expected value of a random variable from a sequence of samples.

$$\hat{X}_{n+1} = \frac{1}{n+1} \sum_{i=1}^{n+1} x_i$$



average of $n+1$ samples

Aside: Online Mean Estimation

- Suppose that we want to incrementally compute the mean of a sequence of numbers (x_1, x_2, x_3, \dots)
 - ▲ E.g. to estimate the expected value of a random variable from a sequence of samples.

$$\hat{X}_{n+1} = \frac{1}{n+1} \sum_{i=1}^{n+1} x_i = \frac{1}{n} \sum_{i=1}^n x_i + \frac{1}{n+1} \left(x_{n+1} - \frac{1}{n} \sum_{i=1}^n x_i \right)$$



average of $n+1$ samples

Aside: Online Mean Estimation

- Suppose that we want to incrementally compute the mean of a sequence of numbers (x_1, x_2, x_3, \dots)
 - ▲ E.g. to estimate the expected value of a random variable from a sequence of samples.

$$\begin{aligned}\hat{X}_{n+1} &= \frac{1}{n+1} \sum_{i=1}^{n+1} x_i = \frac{1}{n} \sum_{i=1}^n x_i + \frac{1}{n+1} \left(x_{n+1} - \frac{1}{n} \sum_{i=1}^n x_i \right) \\ &= \hat{X}_n + \frac{1}{n+1} (x_{n+1} - \hat{X}_n)\end{aligned}$$

average of n+1 samples

learning rate

sample n+1

- Given a new sample x_{n+1} , the new mean is the old estimate (for n samples) plus the weighted difference between the new sample and old estimate

Approach 3: Temporal Difference Learning (TD)

- TD update for transition from s to s' :

$$V^{\pi}(s) \leftarrow V^{\pi}(s) + \alpha(R(s) + \beta V^{\pi}(s') - V^{\pi}(s))$$

updated estimate

learning rate

(noisy) sample of value at s
based on next state s'

- So the update is maintaining a “mean” of the (noisy) value samples
- If the learning rate decreases appropriately with the number of samples (e.g. $1/n$) then the value estimates will converge to true values! (non-trivial)

$$V^{\pi}(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^{\pi}(s')$$

Approach 3: Temporal Difference Learning (TD)

- TD update for transition from s to s' :

$$V^\pi(s) \leftarrow V^\pi(s) + \alpha(R(s) + \beta V^\pi(s') - V^\pi(s))$$

learning rate


(noisy) sample of utility
based on next state

- Intuition about convergence
 - ▶ When V satisfies Bellman constraints then **expected** update is 0.

$$V^\pi(s) = R(s) + \beta \sum_{s'} T(s, a, s') V^\pi(s')$$

- ▶ Can use results from stochastic optimization theory to prove convergence in the limit

The TD learning curve



- **Tradeoff:** requires more training experience (epochs) than ADP but much less computation per epoch
- Choice depends on relative cost of experience vs. computation

Passive RL: Comparisons

- Monte-Carlo Direct Estimation (model free)
 - ▲ Simple to implement
 - ▲ Each update is fast
 - ▲ Does not exploit Bellman constraints
 - ▲ Converges slowly
- Adaptive Dynamic Programming (model based)
 - ▲ Harder to implement
 - ▲ Each update is a full policy evaluation (expensive)
 - ▲ Fully exploits Bellman constraints
 - ▲ Fast convergence (in terms of updates)
- Temporal Difference Learning (model free)
 - ▲ Update speed and implementation similar to direct estimation
 - ▲ Partially exploits Bellman constraints---adjusts state to 'agree' with observed successor
 - Not **all** possible successors as in ADP
 - ▲ Convergence in between direct estimation and ADP

Between ADP and TD

- Moving TD toward ADP
 - ▲ At each step perform TD updates based on observed transition and “imagined” transitions
 - ▲ Imagined transition are generated using estimated model
- The more imagined transitions used, the more like ADP
 - ▲ Making estimate more consistent with next state distribution
 - ▲ Converges in the limit of infinite imagined transitions to ADP
- Trade-off computational and experience efficiency
 - ▲ More imagined transitions require more time per step, but fewer steps of actual experience